

THE BENCHMARK IS NOT THE JOB


Why the AI that tops the leaderboard is often not the one that gets my work done, and how I score them instead.



ONE MODEL, 24 POINTS APART

Same Claude Opus 4.6, same test, three different apps wrapped around it. A leaderboard scores the package, not the model.

Terminal-Bench 2.0 score by agent harness, %




The app built to win the test is not the app I use every day. So I score the two separately.


FOUR LAYERS, ONE SCORE

I rate the setup: a model plus the app I reach it through. The app gets the biggest share.

Weight in the everyday profile, %



Then three gates: privacy, reliability and honesty. Fail one and the score is cut, however good the rest.



NINE THINGS I CHECK

Scored 0 to 4 against things I can see, like whether it asks before deleting.

Gets it done end-to-end	18%
Built-in tools	14%
Connects to your stuff	12%
Remembers context	12%
Recovers from mistakes	12%
Safe control	10%
Speed and reliability	10%
Works everywhere you are	6%
Low learning curve	6%

NOBODY CAN TRAIN ON MY TASKS

Twelve real errands from my own month, run in every app and graded blind with the names hidden.

Usable as-is?	0 · 1 · 2
Follow-ups needed	count
Mistakes I caught	count
Invented sources	count



My time per finished task

My minutes per finished task is the number that matters most. It never maxes out: a better AI just pushes it towards zero.



SOME FLAWS CAN'T BE OUTSCORED

These multiply the score instead of adding to it. Being brilliant elsewhere doesn't buy them back.

GATE	CONCERN	FAIL
Privacy: where my data goes	×0.85	×0.5
Reliability: there when I need it	×0.85	×0.5
Honesty: no invented sources	×0.85	×0.5



IT OUTLASTS THE MODELS

Most leaderboards die when every model scores 95%. This one is built not to.

Skills stay fixed. The tests measuring them get swapped.

A test retires once the leader passes 90%.

Scores are relative: the best setup sets 100.

Every time an AI fails me, that task joins the list.

JUDGE THE ERRAND, NOT THE EXAM

I'll re-run this every time a major model ships, and post the results here.

